

Product Requirement Document [PRD]

Title	Quora – AI Credibility Layer for Trusted Knowledge
Date & Version	4th March, 2026. V1
Product - Point of Contact (POC)	Manish Verma
Design POC	
Tech POC	
Marketing POC	

Objective

1. **Business**:
    a. Improve **30-day retention by 5–10% within 6 months** among high-intent search users by improving answer credibility and information clarity.
    b. Increase **answer engagement depth** by helping users quickly identify the most trustworthy answers.
    c. Reduce **duplicate question fragmentation**, consolidating discussions to improve knowledge discovery.
    d. Strengthen Quora's differentiation as a **trusted knowledge platform** against competitors like Reddit, StackOverflow, and AI search assistants.
2. **Users**:
    a. Help users **quickly identify credible answers** among many responses.
    b. Reduce time spent **scanning low-quality or opinion-heavy answers**.
    c. Provide a **clear summary of the best answers** in long threads.
    d. Improve knowledge discovery by **merging duplicate questions into a single canonical thread**.

How do we measure success?

1. **Goal**: Increase retention among **high-intent search users** by improving answer credibility and reducing information overload.
2. **Success Metrics:**

a. Key Metric: 30-Day Retention Rate (High-Intent Search Users)
b. Secondary Metrics
    i. AI Summary Engagement Rate
    ii. Credibility Label Click Rate
    iii. Average Answers Viewed per Session
    iv. Duplicate Question Reduction Rate

3. **Guardrails:**
    a. Answer Creation Rate
    b. Contributor Retention Rate
    c. Moderation Appeal Rate
    d. Session Load Time

Who are the Users?

1. **Persona:**
    a. **Primary Persona:** Knowledge Seeker
        i. Age: 18–40
        ii. Role: Students, professionals, researchers, learners
        iii. Context: Uses Quora to search for answers, advice, and explanations
        iv. Environment: Users on sessions.
2. **Problems we are solving**
    a. **Core Need**
        i. Users need a way to **identify trustworthy answers quickly** among many responses.
    b. **Problems**
        i. Users struggle to identify the most credible answers in a thread.
        ii. Duplicate questions fragment discussions across multiple threads.
        iii. Long answer lists make it difficult to extract key insights quickly.
        iv. Many answers are opinion-based without credibility indicators.
        v. Users cross-check answers on other platforms, reducing session depth.
    c. **Jobs To Be Done (JTBD)**
        i. **Functional Job**
            1. When I search a question, I want to quickly understand the **best answers**, so that I can get reliable information without reading dozens of responses.
        ii. **Emotional Job**
            1. I want to feel confident that the answer I am reading is **credible and trustworthy**.
        iii. **Social Job**
            1. I want to reference reliable answers confidently when sharing knowledge with others.
3. **How do we know these problems exist?**
    a. **User Interviews (2 conducted)**
        i. Participants reported difficulty identifying the most credible answers.

ii. Users often cross-check answers on Reddit, Google, or AI tools.

b. **Survey (12 responses)**

i. 92% of users reported encountering misleading or incorrect answers.

ii. Many users requested better answer ranking and credibility indicators.

iii. 58% expressed interest in AI-generated summaries of discussions.

c. **Behavioural Signals**

i. Users often leave Quora to verify answers elsewhere.

ii. Duplicate questions create fragmented knowledge threads.

Solution

1. **Brief of the solution:** To address the credibility and fragmentation problem, we will introduce an **AI Credibility Layer** on Quora.

This system will introduce three core improvements:

AI Duplicate Detection

When a user posts a question, AI checks the Quora database for similar questions.

Example detected duplicates:

- Best laptop for coding

• Laptop recommendation for developers

These are automatically merged into **one canonical thread**.

This ensures:

- Knowledge consolidation
- Higher engagement signals
- Reduced duplicate content

AI Credibility Labeling

Each answer will include AI-generated credibility indicators such as:

- Evidence-backed

• Experienced

• Opinion-based

Labels are generated based on:

- sources cited
- author expertise
- engagement signals

This helps users quickly evaluate answer reliability.

AI Summary of Top Answers

AI will generate a summary **only after sufficient community engagement**.

Example:

If a thread contains **20 answers**, AI identifies the **top 10–15% answers by upvotes**.

Example selection:

Top 3 answers → summarised.

Displayed as:

AI Summary – Based on Top Community Answers

Example output:

- MacBook Pro recommended for battery life and performance

• Dell XPS preferred for Linux compatibility

• ThinkPad valued for durability and enterprise reliability

This ensures AI summarises **community-validated knowledge**.

2. **Alternatives Considered (with Prioritisation Metrics):** I evaluated multiple solutions using the **RICE framework** (Reach × Impact × Confidence ÷ Effort).

Solution	Reach (1-3)	Impact (1-3)	Confidence (1-3)	Effort (1-3)	RICE score	Rationale
Credibility Labeling (Fact / Opinion / Experience)	3	3	3	2	13.5	Affects the majority of search-driven users. Directly addresses the trust gap. Strong survey & review backing. Moderate AI + UX effort.
AI Duplicate Question Merge	3	2	3	1	18	High reach across all categories. Reduces noise immediately. Lower engineering effort via NLP clustering. Strong structural improvement.
AI-Generated Thread Summary	3	3	2	3	6	High impact for long threads but AI-dependent and higher compute effort. Needs careful rollout.
Reputation-Based Ranking (Beyond Upvotes)	2	3	2	3	4	Valuable for quality improvement but algorithm redesign is complex and politically sensitive.
Moderation Explanation Assistant	2	2	2	2	4	Improves contributor trust but affects a

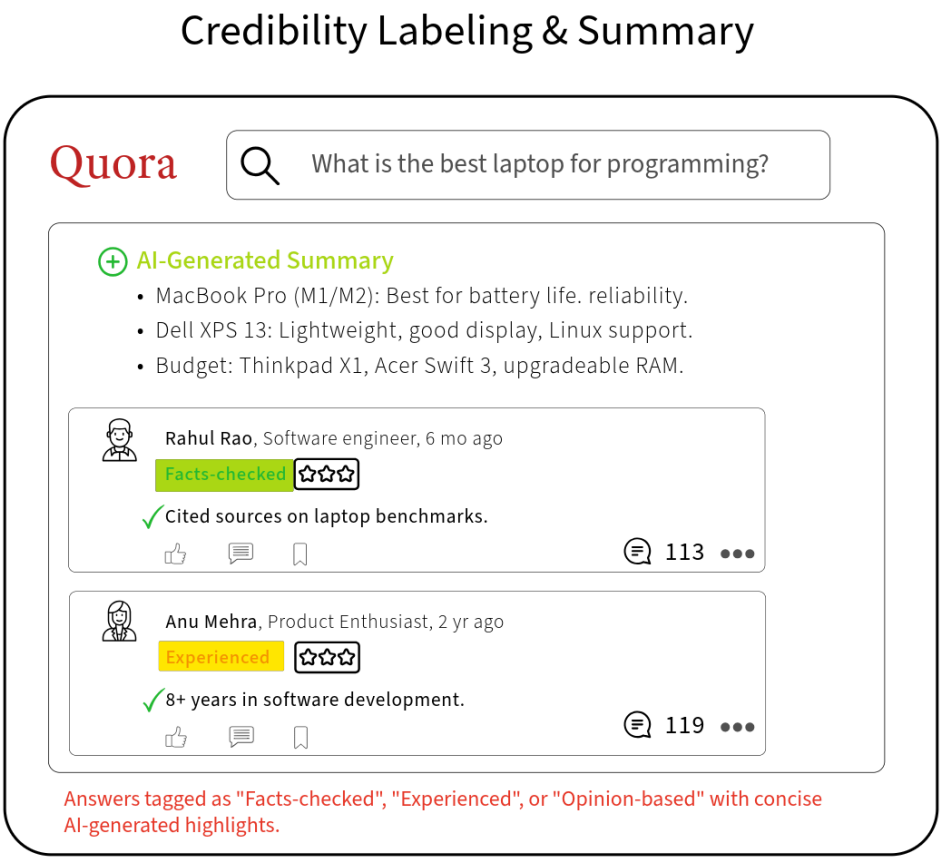
Solution	Reach (1-3)	Impact (1-3)	Confidence (1-3)	Effort (1-3)	RICE score	Rationale
						smaller segment compared to seekers. Moderate build effort.
Canonical Structured Answer (AI Synthesised)	2	3	1	3	1.33	Potentially transformative but assumption-heavy, AI-risk prone, and high engineering effort.

Product flow (Details of the feature/product):

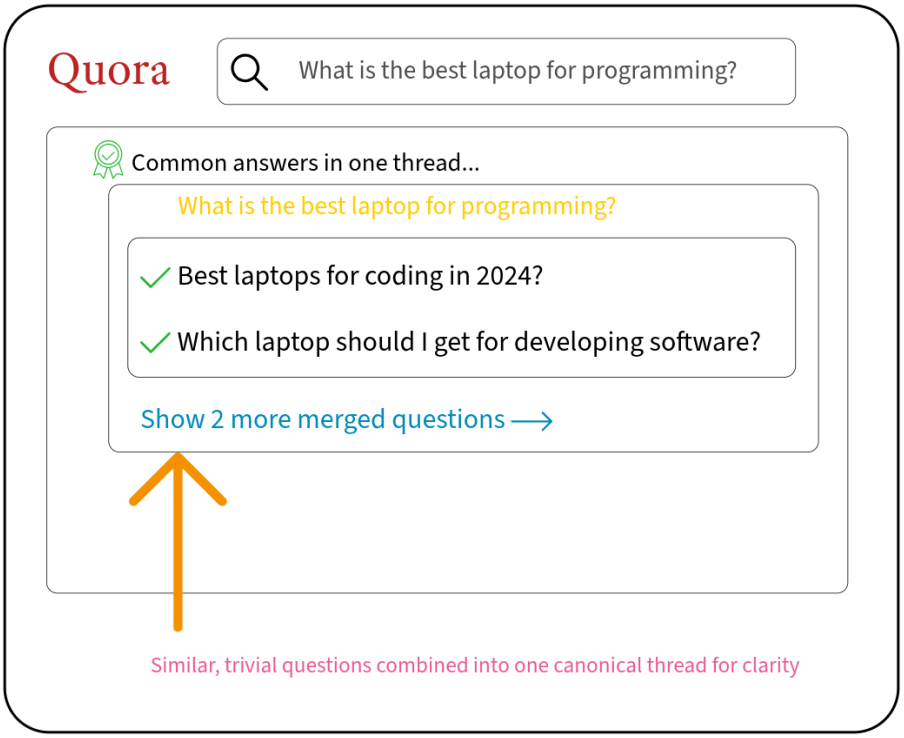
1. **Customer Journey**

   Search Question → Open Thread → AI Detects Duplicate Threads → Read Answers → Community Upvotes → AI Summary Triggered → Users Identify Credible Answers → Continued Engagement

2. **Wireframes and Flow diagrams**



AI duplicate question merge



3. **User Stories**
    a. Story 1 – **Knowledge Seeker:** As a Quora user, I want to quickly identify credible answers so that I can trust the information I read.
    b. Story 2 - **Active Contributor:** As an answer contributor, I want my high-quality answer to be recognised and surfaced in AI summaries.
    c. Story 3 - **New User:** As a new user, I want a quick summary of long answer threads so that I don't need to read dozens of responses.
    d. Story 4 - **Platform Moderator:** As a moderator, I want AI to reduce duplicate questions so that discussions remain organized.
4. **User acceptance criteria**
    a. Duplicate questions are automatically detected and merged.
    b. Credibility labels appear on every answer.
    c. AI summary appears only after engagement threshold.
    d. AI summary uses **top 10–15% upvoted answers**.
    e. Users can still view full answers.
    f. AI labels do not replace community voting.
5. **Edge cases**

Scenario	Expected Behaviour
User posts duplicate question	System suggests existing thread
Thread has very few answers	AI summary not generated

Scenario	Expected Behaviour
Low upvote engagement	AI summary not triggered
Highly controversial topic	AI labels show “Opinion-based”
Spam answers detected	Moderation filters applied

Tentative Timelines:

Task	Date
Leadership approval	01/03/2026
Design ready	14/03/2026
Prototype testing	08/04/2026
Development starts	26/04/2026
Beta Launch	19/10/2026

Dependencies:

1. Open Questions
    a. Should credibility labels be visible to all users or only logged-in users?
    b. Should AI summaries update dynamically as votes change?
    c. Should contributors be able to challenge credibility labels?
2. Infrastructure requirements
    a. AI models for duplicate detection
    b. AI models for credibility labeling
    c. AI summarization pipeline
    d. scalable answer indexing system
    e. ranking system based on engagement signals
3. Budget approvals
    a. AI model infrastructure cost
    b. Engineering development resources
    c. moderation system upgrades
    d. QA testing and rollout monitoring
4. Partner Support (APIs, Partnerships)
    a. Internal search ranking APIs
    b. AI model inference APIs
    c. moderation APIs
5. Internal Dependencies (Tech, Design, Other Teams)
    a. AI Infrastructure Team
    b. Search Ranking Team
    c. Frontend Platform Team
    d. Trust & Safety Team
    e. Product Design Team